

Association mining with PYTHON

Objective: The objective of this exercise is to understand association mining, how frequent itemsets can be extracted by the Apriori algorithm and be able to calculate and interpret association rules in terms of support and confidence.

Material: Lecture notes *"Introduction to Machine Learning and Data Mining"* as well as the files in the exercise 12 folder available from Campusnet.

Discussion forum: You can get help on our online discussion forum:
<https://piazza.com/dtu.dk/spring2023/02450>

Software installation: Extract the Python toolbox from DTU Inside. Start Spyder and add the toolbox directory (`<base-dir>/02450Toolbox_Python/Tools/`) to PYTHONPATH (Tools/PYTHONPATH manager in Spyder). Remember the purpose of the exercises is not to re-write the code from scratch but to work with the scripts provided in the directory `<base-dir>/02450Toolbox_Python/Scripts/` Representation of data in Python:

	Python var.	Type	Size	Description
	X	numpy.array	$N \times M$	Data matrix: The rows correspond to N data objects, each of which contains M attributes.
	attributeNames	list	$M \times 1$	Attribute names: Name (string) for each of the M attributes.
	N	integer	Scalar	Number of data objects.
	M	integer	Scalar	Number of attributes.
Regression	y	numpy.array	$N \times 1$	Dependent variable (output): For each data object, y contains an output value that we wish to predict.
Classification	y	numpy.array	$N \times 1$	Class index: For each data object, y contains a class index, $y_n \in \{0, 1, \dots, C - 1\}$, where C is the total number of classes.
	classNames	list	$C \times 1$	Class names: Name (string) for each of the C classes.
	C	integer	Scalar	Number of classes.
Cross-validation				All variables mentioned above appended with _train or _test represent the corresponding variable for the training or test set.
	*_train	—	—	Training data.
	*_test	—	—	Test data.

12.1 Association Analysis

In this last exercise we will focus on association analysis. Association analysis is widely used in data mining in order to identify important co-occurrence relationships. We will use the following definition of association rule discovery:

Association Rule Discovery. Given a set of transactions T , find all the rules having support $\geq \text{minsup}$ and confidence $\geq \text{minconf}$, where minsup and minconf are the corresponding support and confidence thresholds.

We have summarized the most important terms in table [1](#). We will use the Apriori algorithm to find all itemsets with support greater than $\geq \text{supp}$. The Apriori algorithm is based on the following principle:

Apriori principle. If an itemset is frequent, then all of its subsets must also be frequent.

As a result of the Apriori principle we can start looking at frequent 1-itemsets. The frequent 2-itemsets can then only contain the items in the extracted 1-itemsets and so on and so forth. This greatly reduces the number of itemsets to check to find all frequent itemsets.

Term	Meaning
$I = \{i_1, i_2, \dots, i_d\}$	The set of all items
$T = \{t_1, t_2, \dots, t_N\}$	The set of all transactions
Transaction, t_i	A subset of items: What was bought by a customer
Transaction width	Number of items in transaction
Itemset	A set of items from the set I of all items
k-itemset	An itemset having k items
Support count, $\sigma(X)$	Number of transactions that contain a particular itemset $\sigma(X) = \{t\} $
Association rule	Implication expression of the form $X \leftarrow Y$ where $X \cap Y = \emptyset$
Support, $s(Y \leftarrow X)$	Strength of association rule, $s(Y \leftarrow X) = \frac{\sigma(X \cup Y)}{N} = P(X, Y)$
Confidence, $c(Y \leftarrow X)$	Frequency items in Y appear in transactions containing X, $c(Y \leftarrow X) = \frac{\sigma(X \cup Y)}{\sigma(X)} = \frac{P(X, Y)}{P(X)} = P(Y X)$
Support-based pruning	Pruning strategy based on the Apriori principle (formed by the anti-monotone property)
Anti-monotone property	The support for an itemset never exceeds the support for its subsets (Apriori principle)
F_k	The set of frequent k-itemsets

Table 1: Association mining nomenclature.

12.1.1 In table 2 some of the courses that 6 students completed during their studies are given. Find all itemsets with $\text{supp} \geq 80\%$.

12.1.2 What is the confidence of the rule $02457 \leftarrow 02450$?

	02322	02450	02451	02453	02454	02457	02459	02582
student 1	0	1	0	0	1	1	1	1
student 2	1	1	1	0	0	1	1	1
student 3	0	1	0	1	0	1	0	1
student 4	0	0	1	0	0	1	1	0
student 5	0	1	0	0	0	1	1	0
student 6	0	1	1	0	0	1	1	1

Table 2: Students that upon completing their engineering degree had taken various of the courses 02322, 02450, 02451, 02453, 02454, 02457, 02459 and 02582.

We will use the Apriori algorithm to automatically mine for associations¹. To use the Apriori algorithm, simply install the package `apriori` using `conda install pip` and `pip install apriori` if you are using anaconda (or just `pip install apriori` if you are not using Anaconda).

- 12.1.3 Inspect the file `Data/courses.txt` and the script `ex12_1_3.py`. The script loads the course data file from table 2. Make sure you understand how the data in table 2 is stored in the file and how the script transforms it
- 12.1.4 Inspect and run the script `ex12_1_4.py`. The script transforms the binary matrix, plus label information, into a set of transactions. Make sure you understand how this transformation performs and relate it to the notation of the lecture notes. Finally note how the Apriori algorithm is invoked to find association rules with $\text{supp} \geq 80\%$ and $\text{conf} \geq 100\%$ and print them. Inspect the print command to see how you can access each association rule programmatically. What are the generated association rules?

We will in this last part of the exercise mine for associations in the wine data 1 (<http://archive.ics.uci.edu/ml/datasets/Wine+Quality>) considered in the previous exercises. However, as this data is not binary we will need to convert it to a format suitable for association mining. We will thus binarize the data by dividing each attribute into given percentiles.

- 12.1.5 Inspect and run the script `ex12_1_5.py`. The script load the `Data/wine2.mat` data into Python (for how to load `.mat` files into Python see also exercise

¹A high-performing version of the Apriori algorithm is also available from: <http://www.borgelt.net/apriori.html>.

4.2.1) and divide each of the attributes in the data into percentiles using the function `binarize2`.

The scripts convert the continuous attributes into a one-out-of-K coding based on percentiles. Note how the function also transforms the attribute names. Why do you think we don't just include the 50-100 percentiles of a variable? What are the benefits of including variables corresponding to the 0-50 percentile?

- 12.1.6 Inspect and run the script `ex12_1_6.py` to find association rules in the Wine dataset with $\text{supp} \geq 30\%$ and $\text{conf} \geq 60\%$. Do these association rules make sense?
- 12.1.7 Often we are interested in rules with high confidence. Is it possible for itemsets to have very low support but still have a very high confidence?
- 12.1.8 (optional) Try find associations also in terms of the type of wine by adding two additional columns to the binary data corresponding to `1-y` and `y` (Note this is easiest done by adding new columns to the *X*-matrix and changing the `attributeNames` variable.)

1 Homework problems for this week

Problems

Question 1. Spring 2014 question 11: We consider the twelve costumers given in Table 3. We will consider this data set a market basket problem in which the twelve customers have various combinations of the six items denoted M_H , M_L , P_H , P_L , D_H , D_L . Which one of the proposed solutions below includes *all* the frequent itemsets with support of more than 40 %?

	M_H	M_L	P_H	P_L	D_H	D_L
LIS1	1	0	0	1	1	0
LIS2	1	0	1	0	1	0
LIS3	0	1	0	1	0	1
LIS4	0	1	0	1	0	1
LIS5	1	0	1	0	0	1
LIS6	1	0	1	0	1	0
OPO1	1	0	1	0	0	1
OPO2	0	1	0	1	1	0
OPO3	0	1	1	0	0	1
OPO4	0	1	0	1	0	1
OPO5	0	1	1	0	0	1
OPO6	1	0	1	0	1	0

Table 3: Given are the six first costumers of Lisbon and Oporto including whether these costumers spent more or less than the median consumption of MILK (M_H , M_L), PAPER (P_H , P_L), and DELI (D_H , D_L). Subscript H and L are thus used to respectively denote a relatively high and low level of consumption (i.e., above or below the median consumption).

- A $\{M_L\}, \{M_H\}, \{P_H\}, \{P_L\}, \{D_H\}, \{D_L\}$.
- B $\{M_L\}, \{M_H\}, \{P_H\}, \{P_L\}, \{D_H\}, \{D_L\}, \{M_H, P_H\}$.
- C $\{M_L\}, \{M_H\}, \{P_H\}, \{P_L\}, \{D_H\}, \{D_L\}, \{M_H, P_H\}, \{M_L, D_L\}$.
- D $\{M_L\}, \{M_H\}, \{P_H\}, \{P_L\}, \{D_H\}, \{D_L\}, \{M_H, P_H\}, \{M_H, D_H\}, \{M_L, P_L\}, \{M_L, D_L\}, \{P_L, D_L\}$.
- E Don't know.

Question 2. Fall 2014 question 20: Consider the simple 1-dimensional data set comprised of $N = 7$ observations as shown in table 4. Suppose we wish to apply K-means clustering to the dataset and the $K = 3$ one-dimensional cluster centers are initialized in $\mu_1 = 4$, $\mu_2 = 7$ and $\mu_3 = 14$. After terminating of the K-means clustering algorithm, what are the final (rounded) cluster centers μ_1, μ_2, μ_3 ?

X	3	6	7	9	10	11	14
-----	---	---	---	---	----	----	----

Table 4: Simple 1-dimensional dataset comprised of $N = 7$ observations.

- A $\mu_1 = 3.00, \mu_2 = 8.00, \mu_3 = 12.50$
- B $\mu_1 = 3.00, \mu_2 = 7.33, \mu_3 = 11.67$
- C $\mu_1 = 4.50, \mu_2 = 9.25, \mu_3 = 14.00$
- D $\mu_1 = 5.33, \mu_2 = 10.00, \mu_3 = 14.00$
- E Don't know.

Question 3. Fall 2014 question 11: In table 5 is given the pairwise cityblock distances between 8 observations. A hierarchical clustering is used to cluster these nine observations using *group average* linkage. Which of the dendrograms shown in fig. 1 corresponds to the clustering?

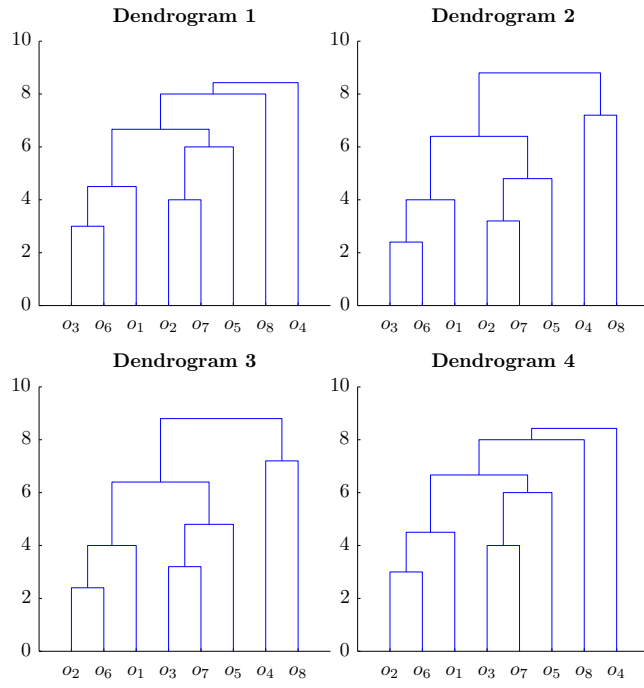


Figure 1: Hierarchical clustering of the 8 observations considered in table 5

	<i>o</i> ₁	<i>o</i> ₂	<i>o</i> ₃	<i>o</i> ₄	<i>o</i> ₅	<i>o</i> ₆	<i>o</i> ₇	<i>o</i> ₈
<i>o</i> ₁	0	4	7	9	5	5	5	6
<i>o</i> ₂	4	0	7	7	7	3	7	8
<i>o</i> ₃	7	7	0	10	6	6	4	9
<i>o</i> ₄	9	7	10	0	8	6	10	9
<i>o</i> ₅	5	7	6	8	0	8	6	7
<i>o</i> ₆	5	3	6	6	8	0	8	11
<i>o</i> ₇	5	7	4	10	6	8	0	7
<i>o</i> ₈	6	8	9	9	7	11	7	0

Table 5: Pairwise Cityblock distance, i.e $d(o_i, o_j) = \|\mathbf{x}_i - \mathbf{x}_j\|_1 = \sum_{k=1}^M |x_{ik} - x_{jk}|$, between 8 observations. Each observation o_i corresponds to a $M = 15$ dimensional binary vector, $x_{ik} \in \{0, 1\}$. The blue observations $\{o_1, o_2, o_3, o_4\}$ belong to class C_1 and the black observations $\{o_5, o_6, o_7, o_8\}$ belong to class C_2 .

- A Dendrogram 1.
- B Dendrogram 2.
- C Dendrogram 3.
- D Dendrogram 4.
- E Don't know.

References

- [1] P. Cortez, A. Cerdeira, F. Almeida, T. Matos, and J. Reis. Modeling wine preferences by data mining from physicochemical properties. *In Decision Support Systems, Elsevier*, 47(4):547–553, 2009.