Best Subset, Forward Stepwise or Lasso? Analysis and Recommendations Based on Extensive Comparisons[1]

Trevor Hastie, Robert Tibshirani and Ryan Tibshirani

Abstract. In exciting recent work, Bertsimas, King and Mazumder (*Ann. Statist.* **44** (2016) 813–852) showed that the classical best subset selection problem in regression modeling can be formulated as a mixed integer optimization (MIO) problem. Using recent advances in MIO algorithms, they demonstrated that best subset selection can now be solved at much larger problem sizes than what was thought possible in the statistics community. They presented empirical comparisons of best subset with other popular variable selection procedures, in particular, the lasso and forward stepwise selection. Surprisingly (to us), their simulations suggested that best subset consistently outperformed both methods in terms of prediction accuracy. Here, we present an expanded set of simulations to shed more light on these comparisons. The summary is roughly as follows:

- neither best subset nor the lasso uniformly dominate the other, with best subset generally performing better in very high signal-to-noise (SNR) ratio regimes, and the lasso better in low SNR regimes;
- for a large proportion of the settings considered, best subset and forward stepwise perform similarly, but in certain cases in the high SNR regime, best subset performs better;
- forward stepwise and best subsets tend to yield sparser models (when tuned on a validation set), especially in the high SNR regime;
- the relaxed lasso (actually, a simplified version of the original relaxed estimator defined in Meinshausen (*Comput. Statist. Data Anal.* **52** (2007) 374–393)) is the overall winner, performing just about as well as the lasso in low SNR scenarios, and nearly as well as best subset in high SNR scenarios.

Key words and phrases: Regression, selection, penalization.

1. INTRODUCTION

Best subset selection, forward stepwise selection and the lasso are popular methods for selection and estimation of the parameters in a linear model. The first two are classical methods in statistics, dating back to at least Beale, Kendall and Mann (1967), Hocking and Leslie (1967) for best subset selection (hereafter "best subset") and Efroymson (1966), Draper and Smith (1966) for forward stepwise selection (hereafter "forward stepwise"); the lasso is (relatively speaking) more recent, due to Tibshirani (1996), Chen, Donoho and Saunders (1998).

Given a response vector $Y \in \mathbf{R}^n$, predictor matrix $X \in \mathbf{R}^{n \times p}$ and a subset size k between 0 and $\min\{n, p\}$, best subset finds the k predictors that produces the best fit

Trevor Hastie is Professor, Departments of Statistics and Biomedical Data Science, Stanford University, Stanford, California 94305, USA (e-mail: hastie@stanford.edu). Robert Tibshirani is Professor, Departments of Statistics and Biomedical Data Science, Stanford University, Stanford, California 94305, USA (e-mail: tibs@stanford.edu). Ryan Tibshirani is Associate Professor, Departments of Statistics and Machine Learning, Carnegie Mellon University, Pittsburgh, Pennsylvania 15217, USA (e-mail: ryantibs@cmu.edu).

in terms of squared error, solving the nonconvex problem

$$\text{(1.1)} \quad \underset{\beta \in \mathbf{R}^p}{\text{minimize}} \|Y - X\beta\|_2^2 \quad \text{subject to} \quad \|\beta\|_0 \le k,$$

where $\|\beta\|_0 = \sum_{i=1}^p 1\{\beta_i \ne 0\}$ is the ℓ_0 norm of β. (Here and throughout, for notational simplicity, we omit the intercept term from the regression model.)

Forward stepwise is less ambitious: starting with the empty model, it iteratively adds the variable that best improves the fit. It hence yields a subset of each size $k = 0, 1, \ldots, \min\{n, p\}$, but none of these are generally globally optimal in the sense of (1.1) (except the first two, and the last if $n > p$). Formally, the procedure starts with an empty active set $A_0 = \{0\}$, and for $k = 1, \ldots, \min\{n, p\}$, selects the variable indexed by

$$\text{(1.2)}$$
$$j_k = \underset{j \notin A_{k-1}}{\text{argmin}} \|Y - P_{A_{k-1} \cup \{j_k\}} Y\|_2^2$$
$$= \underset{j \notin A_{k-1}}{\text{argmax}} \frac{X_j^T P_{A_{k-1}}^\perp Y}{\|P_{A_{k-1}}^\perp X_j\|_2}$$

that leads to the lowest squared error when added to A_{k-1}, or equivalently, such that X_{j_k}, achieves the maximum absolute correlation with Y, after we project out the contributions from $X_{A_{k-1}}$.[1] A note on notation: here, we write $X_S \in \mathbf{R}^{n \times |S|}$ for the submatrix of X whose columns are indexed by a set S (and when $S = \{j\}$, we simply use X_j). We also write P_S for the projection matrix onto the column span of X_S, and $P_S^\perp = I - P_S$ for the projection onto the orthocomplement. At the end of step k of the procedure, the active set is updated, $A_k = A_{k-1} \cup \{j_k\}$, and the forward stepwise estimator of the regression coefficients is defined by the least squares fit onto X_{A_k}. With careful implementation, the computational cost of forward stepwise is equivalent to a single least-squares fit on $\min\{n, p\}$ variables; see Section 2.5.

The lasso solves a convex relaxation of (1.1) where we replace the ℓ_0 norm by the ℓ_1 norm,

$$\text{(1.3)} \quad \underset{\beta \in \mathbf{R}^p}{\text{minimize}} \|Y - X\beta\|_2^2 \quad \text{subject to} \quad \|\beta\|_1 \le t,$$

where $\|\beta\|_1 = \sum_{i=1}^p |\beta_i|$, and $t \ge 0$ is a tuning parameter. By convex duality, the above problem is equivalent to the more common (and more easily solved) penalized form

$$\text{(1.4)} \quad \underset{\beta \in \mathbf{R}^p}{\text{minimize}} \|Y - X\beta\|_2^2 + \lambda \|\beta\|_1,$$

where now $\lambda \ge 0$ is a tuning parameter. This is the form that we focus on in this paper.

The lasso problem (1.4) is convex (and highly structured) and there is by now a sizeable literature in statistics,

machine learning and optimization dedicated to efficient algorithms for this problem. On the other hand, the best subset problem (1.1) is nonconvex and is known to be NP-hard (Natarajan, 1995). The accepted view in statistics for many years has been that this problem is not solveable beyond (say) p in the mid-30s, this view being shaped by the available software for best subset (e.g., in the R language, the `leaps` package implements a branch-and-bound algorithm for best subset of Furnival and Wilson, 1974).

For a much more detailed introduction to best subset, forward stepwise and the lasso, see Chapter 3 of Hastie, Tibshirani and Friedman (2009).

1.1 An Exciting New Development

Recently, Bertsimas, King and Mazumder (2016) presented a mixed integer optimization (MIO) formulation for the best subset problem (1.1). This allows one to use highly optimized MIO solvers, like Gurobi (based on branch-and-cut methods, hybrids of branch-and-bound and cutting plane algorithms), to solve (1.1). Using these MIO solvers, problems with p in the hundreds and even thousands are not out of reach, and this presents us with exciting new ground on which to perform empirical comparisons. Simulation studies in Bertsimas, King and Mazumder (2016) demonstrated that best subset generally gives superior prediction accuracy compared to forward stepwise and the lasso, over a variety of problem setups.

In what follows, we replicate and expand these simulations to shed more light on such comparisons. For convenience, we created an R package `bestsubset` for optimizing the best subset problem using the Gurobi MIO solver (after this problem has been translated into a mixed integer quadratic program as in Bertsimas, King and Mazumder, 2016). This package, as well as R code for reproducing all of the results in this paper, are available at https://github.com/ryantibs/best-subset/.

1.2 Other Estimators

Many other sparse estimators for regression could be considered, for example, ℓ_1-penalized alternatives to the lasso, like the Dantzig selector (Candes and Tao, 2007) and square-root lasso (Belloni, Chernozhukov and Wang, 2011); greedy alternatives to the forward stepwise algorithm, like matching pursuit (Zhang, 1993) and orthogonal matching pursuit (Davis, Mallat and Zhang, 1994); nonconvex-penalized methods, such as SCAD (Fan and Li, 2001), MC+ (Zhang, 2010) and SparseNet (Mazumder, Friedman and Hastie, 2011); hybrid lasso/stepwise approaches like FLASH (Radchenko and James, 2011); and many others.

It would be interesting to include all of these estimators in our comparisons, though that would make for a huge simulation suite and would dilute the comparisons between best subset, forward stepwise, and the lasso that

---

[1]Other ways of defining the variable j_k that "best improves the fit" are possible, but the entry criterion is (1.2) is the standard one in statistics.

we would like to highlight. Roughly speaking, we would expect the Dantzig selector and square-root lasso to perform similarly to the lasso; the matching pursuit variants to perform similarly to forward stepwise; and the nonconvex-penalized methods to perform somewhere in between the lasso and best subset. (It is worth noting that our R package is structured in such a way to make further simulations and comparisons straightforward. We invite interested readers to use it to perform comparisons to other methods.) Indeed, in the Supplementary Appendix (Hastie, Tibshirani and Tibshirani, 2020) we include comparisons with three additional methods.

1.3 What This Paper Is About

We need to make clear the focus of this paper. This paper is *not* about:

- What is the best prediction algorithm?
- What is the best variable selector?
- Empirically validating theory for ℓ_0 and ℓ_1 penalties.

Rather, this paper is about:

> *The relative merits of the three (arguably) most canonical forms for sparse estimation in a linear model: ℓ_0, ℓ_1 and forward stepwise selection.*

2. PRELIMINARY DISCUSSION

2.1 Is Best Subset the Holy Grail?

Various researchers throughout the years have viewed best subset as the "holy grail" of estimators for sparse modeling in regression, suggesting (perhaps implicitly) that it should be used whenever possible, and that other methods for sparse regression—such as forward stepwise and the lasso—should be seen as approximations or heuristics, used only out of necessity when best subset is not computable. However, as we will demonstrate in the simulations that follow, this is not the case. Different procedures have different operating characteristics, that is, give rise to different bias-variance tradeoffs as we vary their respective tuning parameters. In fact, depending on the problem setting, the bias-variance tradeoff provided by best subset may be more or less useful than the tradeoff provided by the lasso.

As a brief interlude, let us inspect the "noiseless" versions of the best subset and lasso optimization problems. Suppose we observe n examples of a linear system with $p > n$ unknown parameters $Y = X\beta$. This system is underdetermined, and we wish to learn the parameter β. Since there are in general infinitely many solutions, one approach is to seek the sparsest:

$$(2.1) \qquad \underset{\beta \in \mathbf{R}^p}{\text{minimize}} \|\beta\|_0 \quad \text{subject to} \quad X\beta = Y.$$

Since this problem is nonconvex and in general NP hard, we might solve instead the ℓ_1 relaxation:

$$(2.2) \qquad \underset{\beta \in \mathbf{R}^p}{\text{minimize}} \|\beta\|_1 \quad \text{subject to} \quad X\beta = Y.$$

If our goal is truly to seek the sparsest solution to $Y = X\beta$, then problem (2.1) produces it, and we may rightly view problem (2.2) as a heuristic. Indeed, much of the literature on compressed sensing (in which (2.1) and (2.2) have been intensely studied) uses this language. However, one must not be careful to blindly carry this mindset over to the "noisy" setting, which is the traditional and most practical setting for statistical estimation, and that studied in this paper. When there is observational noise, the bias and variance of an estimator play big roles (the bias-variance tradeoff appears), and whether ℓ_0- or ℓ_1-regularization delivers a "better" estimator is a subtle question.

Generally speaking, the lasso and best subset differ in terms of their "aggressiveness" in selecting and estimating the coefficients in a linear model, with the lasso being less aggressive than best subset; forward stepwise lands somewhere in the middle. There are various ways to make this vague but intuitive comparison more explicit. For example:

- forward stepwise can be seen as a "locally optimal" version of best subset, updating the active set by one variable at each step, instead of reoptimizing over all possible subsets of a given size; in turn, the lasso can be seen as a more "democratic" version of forward stepwise, updating the coefficients so as to maintain equal absolute correlation of all active variables with the residual (Efron et al., 2004);
- the lasso applies shrinkage to its nonzero estimated coefficients (e.g., see (2.5)) but forward stepwise and best subset do not, and simply perform least squares on their respective active sets;
- thanks to such shrinkage, the fitted values from the lasso (for any fixed $\lambda \geq 0$) are continuous functions of y (Zou, Hastie and Tibshirani, 2007, Tibshirani and Taylor, 2012), whereas the fitted values from forward stepwise and best subset (for fixed $k \geq 1$) jump discontinuously as y moves across a decision boundary for the active set;
- again thanks to shrinkage, the effective degrees of freedom of the lasso (at any fixed $\lambda \geq 0$) is equal to the expected number of selected variables (Zou, Hastie and Tibshirani, 2007, Tibshirani and Taylor, 2012), whereas the degrees of freedom of both forward stepwise and best subset can greatly exceed k at any given step $k \geq 1$ (Kaufman and Rosset, 2014, Janson, Fithian and Hastie, 2015). Effective degrees of freedom is a useful measure of complexity, especially for models that are fit adaptively. Figure 1 uses effective degrees of freedom to contrast the aggressiveness of the three methods.

FIG. 1.   *Effective degrees of freedom $\sum_{i=1}^{n} \text{Cov}(Y_i, \hat{Y}_i)/\sigma^2$ for the lasso, forward stepwise and best subset, in a problem setup with $n = 70$ and $p = 30$ (computed via Monte Carlo evaluation of the covariance with 500 replications). The setup had an SNR of 0.7, predictor correlation level of 0.35 and the coefficients followed the beta-type 2 pattern with $s = 5$; see Section 3.1 for details. Note that the lasso degrees of freedom equals the (expected) number of nonzero coefficients, whereas that of forward stepwise and best subset exceeds the number of nonzero coefficients.*

When the signal-to-noise ratio (SNR) is low, and also depending on other factors like the correlations between predictor variables, the more aggressive best subset and forward stepwise methods can already have quite high variance at the start of their model paths (i.e., for small step numbers k). Even after optimizing over the tuning parameter k (using say, an external validation set or an oracle which reveals the true risk), we can arrive at an estimator with unwanted variance and worse accuracy than a properly-tuned lasso estimator. On the other hand, for high SNR values, and other configurations for the correlations between predictors, etc., the story can be completely flipped and the shrinkage applied by the lasso estimator can result in unwanted bias and worse accuracy than best subset and forward stepwise. See Figure 2 for empirical evidence.

This is a simple point, but is worth emphasizing:

> *Different procedures bring us from the high-bias to the high-variance ends of the tradeoff along different model paths; and these paths are affected by aspects of the problem setting, like the SNR and predictor correlations, in different ways. For some classes of problems, some procedures admit more fruitful paths, and for other classes, other procedures admit more fruitful paths. For example, neither*

> *best subset nor the lasso dominates the other, across all problem settings.*

2.2 What Is a Realistic Signal-to-Noise Ratio?

In their simulation studies, Bertsimas, King and Mazumder (2016) considered SNRs in the range of about 2 to 8 in their low-dimensional cases, and about 3 to 10 in their high-dimensional cases. Is this a realistic range that one encounters in practice? In our view, the proportion of variance explained (PVE) can help to answer this question.

Let $(x_0, y_0) \in \mathbf{R}^p \times \mathbf{R}$ be a pair of predictor and response variables, and define $f(x_0) = \mathbf{E}(y_0|x_0)$ and $\epsilon_0 = y_0 - f(x_0)$, so that we may express the relationship between x_0, y_0 as

$$y_0 = f(x_0) + \epsilon_0.$$

The signal-to-noise ratio (SNR) in this model is defined as

$$\text{SNR} = \frac{\text{Var}(f(x_0))}{\text{Var}(\epsilon_0)}.$$

The proportion of variance explained (PVE) by a candidate prediction function g is defined as

$$(2.3) \qquad \text{PVE}(g) = 1 - \frac{\mathbf{E}(y_0 - g(x_0))^2}{\text{Var}(y_0)}$$

FIG. 2. *Relative risk (risk divided by null risk) for the lasso, forward stepwise and best subset, for two different setups. The results were averaged over 20 repetitions, and the bars denote one standard error. The setup for the left panel is identical to that used in Figure 1. The setup for the right panel used an SNR (signal-to-noise ratio) of 2 and predictor correlation level 0. Note that in the left panel, the lasso is more accurate than forward stepwise selection and best subset selection (when all methods are optimally-tuned), and in the right panel, the opposite is true.*

(it is 1 minus the proportion of unexplained variance). Of course, this is maximized when we take g to be the mean function f itself, in which case

$$(2.4) \qquad \text{PVE}(f) = 1 - \frac{\text{Var}(\epsilon_0)}{\text{Var}(y_0)} = \frac{\text{SNR}}{1 + \text{SNR}}.$$

In the second equality, we have assumed independence of x_0 and ϵ_0, so $\text{Var}(y_0) = \text{Var}(f(x_0)) + \text{Var}(\epsilon_0)$. As the optimal prediction function is f, it sets the gold-standard of $\text{SNR}/(1 + \text{SNR})$ for the PVE, so we should always expect to see the attained PVE be less than $\text{SNR}/(1 + \text{SNR})$ and greater than 0 (otherwise we could simply replace our prediction function by $g = 0$).

We illustrate using a simulation with $n = 200$ and $p = 100$. The predictor correlation level was set to zero and the coefficients followed the beta-type 2 pattern with $s = 5$;


FIG. 3. *PVE (proportion of variance explained) of the lasso in a simulation setup with $n = 200$ and $p = 100$, as the SNR varies from 0.05 to 6 (more details are provided in the text). The red curve is the population PVE, the maximum achievable PVE at any given SNR value. We see that SNRs above 2 give PVEs roughly above 0.6, which seems to us to be rare in many practical applications.*

see Section 3.1 for details. We varied the SNR in the simulation from 0.05 to 6 in 20 equally spaced values. We computed the lasso over 50 values of the tuning parameter λ, and selected the tuning parameter by optimizing prediction error on a separate validation set of size n. Figure 3 shows the PVE of the tuned lasso estimator, averaged over 20 repetitions from this simulation setup. Also shown is the maximal population PVE (2.4). We see that a SNR of 1.0 corresponds to a PVE of about 0.45 (with a maximum of 0.5), while a SNR as low as 0.25 yields a PVE of 0.1 (with a maximum of 0.2). In our experience, a PVE of 0.5 is rare for noisy observational data, and 0.2 may be more typical. A PVE of 0.86, corresponding to a SNR of 6, is unheard of. With financial returns data, explaining even 2% of the variance (PVE of 0.02) would be considered huge, and the corresponding prediction function could lead to considerable profits if used in a trading scheme. Therefore, based on these observations, we examine a wider range of SNRs in our simulations, compared to the SNRs studied in Bertsimas, King and Mazumder (2016).

2.3 A (Simplified) Relaxed Lasso

In addition to the lasso estimator, we consider a simplified version of the relaxed lasso estimator as originally defined by Meinshausen (2007). Let $\hat{\beta}^{\text{lasso}}(\lambda)$ denote the solution in problem (1.4), that is, the lasso estimator at the tuning parameter value $\lambda \geq 0$. Let A_λ denote its active set, and let $\hat{\beta}^{\text{LS}}_{A_\lambda}$ denote the least squares coefficients obtained by regressing of Y on X_{A_λ}, the submatrix of active predictors. Finally, let $\hat{\beta}^{\text{LS}}(\lambda)$ be the full-sized (p-dimensional) version of the least squares coefficients, padded with zeros to match the zeros of the lasso solution. We consider the estimator $\hat{\beta}^{\text{relax}}(\lambda, \gamma)$

$$(2.5) \qquad \hat{\beta}^{\text{relax}}(\lambda, \gamma) = \gamma \hat{\beta}^{\text{lasso}}(\lambda) + (1 - \gamma)\hat{\beta}^{\text{LS}}(\lambda)$$

FIG. 4.   *Degrees of freedom for the lasso, forward stepwise, best subset and the relaxed lasso with $\gamma = 0.5$ and $\gamma = 0$. The problem setup is the same as that in the left panel of Figure 1. Note that the relaxed lasso has an inflated degrees of freedom compared to the lasso and generally has a larger degrees of freedom than the expected number of nonzero coefficients. But, even when $\gamma = 0$, its degrees of freedom is smaller than that of forward stepwise and best subset throughout their model paths.*

with respect to the pair of tuning parameter values $\lambda \geq 0$ and $\gamma \in [0, 1]$. Recall (Tibshirani, 2013) that when the columns of X are in general position (a weak condition occurring almost surely for continuously distributed predictors, regardless of n, p), it holds that:

- the lasso solution is unique;
- the submatrix X_{A_λ} of active predictors has full column rank, thus

$$\hat{\beta}_{A_\lambda}^{\text{LS}} = (X_{A_\lambda}^T X_{A_\lambda})^{-1} X_{A_\lambda}^T Y$$

  is well-defined;
- the lasso solution can be written (over its active set) as

$$\hat{\beta}_{A_\lambda}^{\text{lasso}}(\lambda) = (X_{A_\lambda}^T X_{A_\lambda})^{-1}(X_{A_\lambda}^T Y - \lambda s),$$

where $s \in \{-1, 1\}^{|A_\lambda|}$ contains the signs of the active lasso coefficients.

Thus, under the general position assumption on X, the simplified relaxed lasso can be rewritten as

$$\hat{\beta}_{A_\lambda}^{\text{relax}}(\lambda, \gamma) = (X_{A_\lambda}^T X_{A_\lambda})^{-1} X_{A_\lambda}^T Y$$

(2.6)
$$- \gamma \lambda (X_{A_\lambda}^T X_{A_\lambda})^{-1} s,$$

$$\hat{\beta}_{-A_\lambda}^{\text{relax}}(\lambda, \gamma) = 0,$$

so we see that $\gamma \in [0, 1]$ acts as a multiplicative factor applied directly to the "extra" shrinkage term apparent in

the lasso coefficients. Henceforth, we will drop the word "simplified" and will just refer to this estimator as the relaxed lasso.

The relaxed lasso tries to undo the shrinkage inherent in the lasso estimator, to a varying degree, depending on γ. In this sense, we would expect it to be more aggressive than the lasso, and have a larger effective degrees of freedom. However, even in its most aggressive mode, $\gamma = 0$, it is typically less aggressive than both forward stepwise and best subset, in that it often has a smaller degrees of freedom than these two. See Figure 4 for an example.

2.4 Some Alternatives to the Bertsimas, King and Mazumder (2016) Approach for Subset Selection

After the completion of the first version of this paper, two refinements to the MIO approach were proposed in Bertsimas and Van Parys (2020) and Hazimeh and Mazumder (2018). We tried two methods from the latter paper ("L0Learn 1" and "L0Learn 2") and also "SparseNet" (Mazumder, Friedman and Hastie, 2011) in our comparisons. SparseNet builds a family of regularization paths spanning ℓ_1 and ℓ_0, using a mixture penalty inspired by the elastic net. L0Learn1 and L0Learn2 offered little improvement over best subset; SparseNet performed reasonably well, but was generally outperformed by the relaxed lasso family. Hence we do not include them here

but provide full results in the Supplementary Appendix (Hastie, Tibshirani and Tibshirani, 2020).

2.5 Brief Discussion of Computational Costs

Computation of the lasso solution in (1.4) has been a popular topic of research, and there are by now many efficient lasso algorithms. In our simulations, we use coordinate descent with warm starts over a sequence of tuning parameter values $\lambda_1 > \cdots > \lambda_m > 0$, as implemented in the glmnet R package (Friedman et al., 2007, Friedman, Hastie and Tibshirani, 2010). The base code for this is written in Fortran, and warm starts—plus additional tricks like active set optimization and screening rules (Tibshirani et al., 2012)—make this implementation highly efficient. For example, for a problem with $n = 500$ observations and $p = 100$ variables, glmnet delivers the lasso solutions across 100 values of λ in less than 0.01 seconds, on a standard laptop computer. The relaxed lasso in (2.5) comes at only a slight increase in computational cost, since we must only additionally compute the least squares coefficients on each active set. We provide an implementation in the bestsubset R package accompanying this paper, which just uses an R wrapper around glmnet. A future version of glmnet will include the relaxed lasso. For the same example with $n = 500$ and $p = 100$, computing the relaxed lasso path over 100 values of λ and 10 values of γ again took less than 0.01 seconds.

For forward stepwise, we implemented our own version in the bestsubset R package. The core matrix manipulations for this method are written in C, and the rest is written in R. The forward stepwise path is highly structured and this greatly aids its computation: at step k, we have $k - 1$ active variables included in the model, and we seek the variable among the remaining $p - k + 1$ that— once orthogonalized with respect to the current active set of variables—achieves the greatest absolute correlation with Y, as in (1.2). Suppose that we have maintained a QR decomposition of the active submatrix $X_{A_{k-1}}$ of predictors, as well as the orthogonalization of the remaining $p - k - 1$ predictors with respect to $X_{A_{k-1}}$. We can compute the necessary correlations in $O(n(p - k + 1))$ operations, update the QR factorization of X_{A_k} in constant time, and orthogonalize the remaining predictors with respect to the one just included in $O(n(p - k))$ operations (refer to the modified Gram–Schmidt algorithm in Golub and Van Loan, 1996). Hence, the forward stepwise path can be seen as a certain guided QR decomposition for computing the least squares coefficients on all p variables (or, on some subset of n variables when $p > n$). For the same example with $n = 500$ and $p = 100$, our implementation computes the forward stepwise path in less than 0.5 seconds.

Best subset (1.1) is the most computationally challenging, by a large margin. Bertsimas, King and Mazumder (2016) describe two reformulations of (1.1) as a mixed integer quadratic program, one that is preferred when $n \geq p$, and the other when $p > n$, and recommend using the Gurobi commercial MIO solver (which is free for academic use). They also describe a proximal gradient descent method for computing approximate solutions in (1.1), and recommend using the best output from this algorithm over many randomly-initialized runs to warm start the Gurobi solver; see Bertsimas, King and Mazumder (2016) for details. We have implemented the method of these authors, which transforms the best subset problem into one of two MIO formulations depending on the relative sizes of n and p, uses proximal gradient to compute a warm start, and then calls Gurobi through its R interface—in our accompanying R package bestsubset.

Gurobi uses branch-and-cut techniques (a combination of branch-and-bound and cutting plane methods), along with many other sophisticated optimization tools, for MIO problems. Compared to the pure branch-and-bound method from the leaps R package, its speed can be impressive: for example, in one run with $n = 500$ and $p = 100$, it returned the best subset solution of size $k = 8$ in about 3 minutes (brute-force search for this problem would need to have looked at about 186 billion candidates). But for most problems of this size ($n = 500$ and $p = 100$) it has been our experience that Gurobi typically requires 1 hour or longer to complete its optimization. It can often be the case that Gurobi has found the solution in less than 3 minutes, though it takes much longer to certify its optimality. For our simulations in the next section, we used a time limit of 30 minutes for Gurobi to optimize the best subset problem (1.1) at any particular value of the subset size k (once the time limit has been reached, the solver returns its best iterate). For more discussion on this choice and its implications, see Section 3.2. We note that this corresponds to a computational cost for "regular" practical usage of 30 minutes per value of k: if we wanted to use 10-fold cross-validation to choose between the subset sizes $k = 0, \ldots, 50$, then we are facing 250 hours (> 10 days) of computation time.

3. SIMULATIONS

3.1 Setup

We present simulations, basically following the simulation setup of Bertsimas, King and Mazumder (2016), except that we consider a wider range of SNR values. Given n, p (problem dimensions), s (sparsity level), beta-type (pattern of sparsity), ρ (predictor correlation level) and ν (SNR level), our process can be described as follows:

(i) we define coefficients $\beta_0 \in \mathbf{R}^p$ according to s and the beta-type, as described below;

(ii) we draw the rows of the predictor matrix $X \in \mathbf{R}^{n \times p}$ i.i.d. from $N_p(0, \Sigma)$, where $\Sigma \in \mathbf{R}^{p \times p}$ has entry (i, j) equal to $\rho^{|i-j|}$;

(iii) we draw the response vector $Y \in \mathbf{R}^n$ from $N_n(X\beta_0, \sigma^2 I)$, with σ^2 defined to meet the desired SNR level, that is, $\sigma^2 = \beta_0^T \Sigma \beta_0 / \nu$;

(iv) we run the lasso, relaxed lasso, forward stepwise, and best subset on the data X, Y, each over a wide range of tuning parameter values; for each method, we choose the tuning parameter by minimizing prediction error on a validation set $\widetilde{X} \in \mathbf{R}^{n \times p}$, $\widetilde{Y} \in \mathbf{R}^n$ that is generated independently of and identically to X, Y, as in steps (ii)–(iii) above;

(v) we record several metrics of interest, as specified below;

(vi) we repeat steps (ii)–(v) a total of 10 times, and average the results.

Below we describe some aspects of the simulation process in more detail.

Coefficients. We considered four settings for the coefficients $\beta_0 \in \mathbf{R}^p$:

- *beta-type 1*: β_0 has s components equal to 1, occurring at (roughly) equally-spaced indices between 1 and p, and the rest equal to 0;
- *beta-type 2*: β_0 has its first s components equal to 1, and the rest equal to 0;
- *beta-type 3*: β_0 has its first s components taking nonzero values equally-spaced between 10 and 0.5, and the rest equal to 0;
- *beta-type 5*: β_0 has its first s components equal to 1, and the rest decaying exponentially to 0, specifically, $\beta_{0i} = 0.5^{i-s}$, for $i = s+1, \ldots, p$.

The first three types were studied in Bertsimas, King and Mazumder (2016). They also defined a fourth type that we did not include here, as we found it yielded basically the same results as beta-type 3. The last type above is new: we included it to investigate the effects of weak sparsity and call it beta-type 5, to avoid confusion.

Evaluation metrics. Let $x_0 \in \mathbf{R}^p$ denote test predictor values drawn from $N_p(0, \Sigma)$ (as in the rows of the training predictor matrix X) and let $y_0 \in \mathbf{R}$ denote its associated response value drawn from $N(x_0^T \beta_0, \sigma^2)$. Also, let $\hat{\beta}$ denote estimated coefficients from one of the regression procedures. We considered the following evaluation metrics:

- *Relative risk*: this is the accuracy metric studied in Bertsimas, King and Mazumder (2016)[2], defined as

$$\mathrm{RR}(\hat{\beta}) = \frac{\mathbf{E}(x_0^T \hat{\beta} - x_0^T \beta_0)^2}{\mathbf{E}(x_0^T \beta_0)^2}$$

$$= \frac{(\hat{\beta} - \beta_0)^T \Sigma (\hat{\beta} - \beta_0)}{\beta_0^T \Sigma \beta_0}.$$

The expectations here and below are taken over the test point (x_0, y_0), with all training data and validation data and thus $\hat{\beta}$) held fixed. A perfect score is 0 (if $\hat{\beta} = \beta_0$) and the null score is 1 (if $\hat{\beta} = 0$).

- *Relative test error*: this measures the expected test error relative to the Bayes error rate,

$$\mathrm{RTE}(\hat{\beta}) = \frac{\mathbf{E}(y_0 - x_0^T \hat{\beta})^2}{\sigma^2}$$

$$= \frac{(\hat{\beta} - \beta_0)^T \Sigma (\hat{\beta} - \beta_0) + \sigma^2}{\sigma^2}.$$

A perfect score is 1 and the null score is

$$(\beta_0^T \Sigma \beta_0 + \sigma^2)/\sigma^2 = \mathrm{SNR} + 1.$$

- *Proportion of variance explained*: as defined in Section 2.2, this is,

$$\mathrm{PVE}(\hat{\beta}) = 1 - \frac{\mathbf{E}(y_0 - x_0^T \hat{\beta})^2}{\mathrm{Var}(y_0)}$$

$$= 1 - \frac{(\hat{\beta} - \beta_0)^T \Sigma (\hat{\beta} - \beta_0) + \sigma^2}{\beta_0^T \Sigma \beta_0 + \sigma^2}.$$

A perfect score is $\mathrm{SNR}/(1 + \mathrm{SNR})$ and the null score is 0.

- *Number of nonzeros*: unlike the previous two metrics which measure predictive accuracy, this metric simply records the number of nonzero estimated coefficients, $\|\hat{\beta}\|_0 = \sum_{i=1}^p 1\{\hat{\beta}_i \neq 0\}$, which we compare with the true value.

- *F-score*: this measures the accuracy of the support recovery:

$$\mathrm{F\text{-}score} = \left(\frac{\mathrm{recall}^{-1} + \mathrm{precision}^{-1}}{2} \right)^{-1},$$

where in statistical terminology, recall is the sensitivity (true positive rate) and precision is the positive predictive value. An F-score of 1 predicts perfectly.

A few notes are in order. For brevity, we do not include relative risk (RR) in any of our plots here, though

---

[2] Actually, these authors used an "in-sample" version of this metric defined as $\|X\hat{\beta} - X\beta_0\|_2^2 / \|X\hat{\beta}\|_2^2$, whereas our definition is "out-of-sample", with an expectation over the new test predictor value x_0 taking the place of the sample average over the training values x_i, $i = 1, \ldots, n$.

it is included in the Supplementary Material (Hastie, Tibshirani and Tibshirani, 2020). Also, we mention that the prediction-based metrics above (RR, RTE, PVE) do not depend on the assumption of linearity of $\mathbf{E}(y|x)$, which reflects the broader practical relevancy of these metrics (compared to F-score).

Configurations. We considered the following four problem settings:

Setting	n	p	s
Low	100	10	5
Medium	500	100	5
High-5	50	1000	5
High-10	100	1000	10

In each setting, we considered ten values for the SNR ranging from 0.05 to 6 on a log scale, namely

SNR	0.05	0.09	0.14	0.25	0.42	0.71	1.22	2.07	3.52	6.00
PVE	0.05	0.08	0.12	0.20	0.30	0.42	0.55	0.67	0.78	0.86

(Note that although these values are focused on the smaller SNR values, they are more evenly spread out on the PVE scale, which we have also provided for convenience.) In each setting, we also considered three values for the predictor correlation level ρ, namely 0, 0.35 and 0.7.

Tuning of procedures. In the low setting, the lasso was tuned over 50 values of λ ranging from $\lambda_{\max} = \|X^T Y\|_\infty$ to a small fraction of $\lambda_{\max}$ on a log scale, as per the default in glmnet, and the relaxed lasso was tuned over the same 50 values of λ, and 10 values of γ equally spaced from 1 to 0 (hence a total of 500 tuning parameter values). Also in the low setting, forward stepwise and best subset were tuned over steps $k = 0, \ldots, 10$. In all other problem settings (medium, high-5, and high-10), the lasso was tuned over 100 values of λ, the relaxed lasso was tuned over the same 100 values of λ and 10 values of γ (hence 1000 tuning parameter values total), and forward stepwise and best subset were tuned over steps $k = 0, \ldots, 50$. In all cases, tuning was performed by minimizing prediction error on an external validation set of size n, which we note approximately matches the precision of leave-one-out cross-validation.

3.2 Time Budget for Gurobi

As mentioned in Section 2.5, for each problem instance and subset size k, we used a time limit of 30 minutes for Gurobi to optimize the best subset problem. In comparison, Bertsimas, King and Mazumder (2016) used 15 minutes (per problem per k) for problems with $p = 100$ as in our medium setup, and 66 minutes (per problem per k) for problems with $p \geq 1000$ as in our high-5 and high-10 setups. Their simulations however were not as extensive, as they looked at fewer combinations of beta-types, SNR levels and correlation levels.

Time in seconds for one path of solutions for best subset (BS), forward stepwise (FS), the lasso and relaxed lasso (RLasso). The times were averaged over 10 repetitions, and all SNR and predictor correlation levels in the given setting

Setting	n	p	s	BS	FS	Lasso	RLasso
Low	100	10	5	0.313	0.003	0.002	0.002
Medium	500	100	5	76.8 hr	0.890	0.013	0.154
High-5	50	100	5	44.2 hr	0.123	0.014	0.159
High-10	100	1000	10	61.7 hr	0.254	0.024	0.158

The MIO solver in the medium setting will often arrive at the best subset solution in less than 30 minutes, but it can take much longer to certify its optimality[3] (usually over 1 hour, in absence of extra speedup tricks as described in Bertsimas, King and Mazumder, 2016). For practical reasons, we have kept the 30 minute budget per problem instance per subset size. Note that this amounts to 1500 minutes per path of 50 solutions, 15,000 minutes or 250 hours per set of 10 repetitions, and in total 7500 hours or 312.5 days for any given setting, once we go through the 10 SNR levels and 3 correlation levels. Fortunately, we had access to a large cluster where we could reduce this time by a factor of about 50.

3.3 Results: Computation Time

In Table 1, we report the time in seconds taken by each method to compute one path of solutions, averaged over 10 repetitions and all SNR and predictor correlation levels in the given setting. All timings were recorded on a Linux cluster. As explained above, the lasso path consisted of 50 tuning parameter values in the low setting and 100 in all other settings, the relaxed lasso path consisted of 500 tuning parameter values in the low setting and 1000 in all other settings, and the forward stepwise and best subset paths each consisted of $\min\{p, 50\}$ tuning parameter values.

We can see that the lasso and relaxed lasso are very fast, requiring less than 25 milliseconds in every case. forward stepwise is also fast, though not quite as fast as the lasso (some of the differences here might be due to the fact that our forward stepwise algorithm is implemented partly in R). Moreover, it should be noted that when n and p is large, and one wants to explore models with a sizeable number of variables (we limited our search to models of size 50), forward stepwise has to plod through its path one variable at a time, but the lasso can make jumps over subset sizes bigger than one by varying λ and leveraging warm starts.

---

[3]Gurobi constructs a sequence of lower and upper bounds on the criterion in (1.1); typically the lower bounds come from convex relaxations and the upper bounds from the current iterates, and it is the lower bounds that take so long to converge.

Low setting: $n = 100$, $p = 10$, $s = 5$
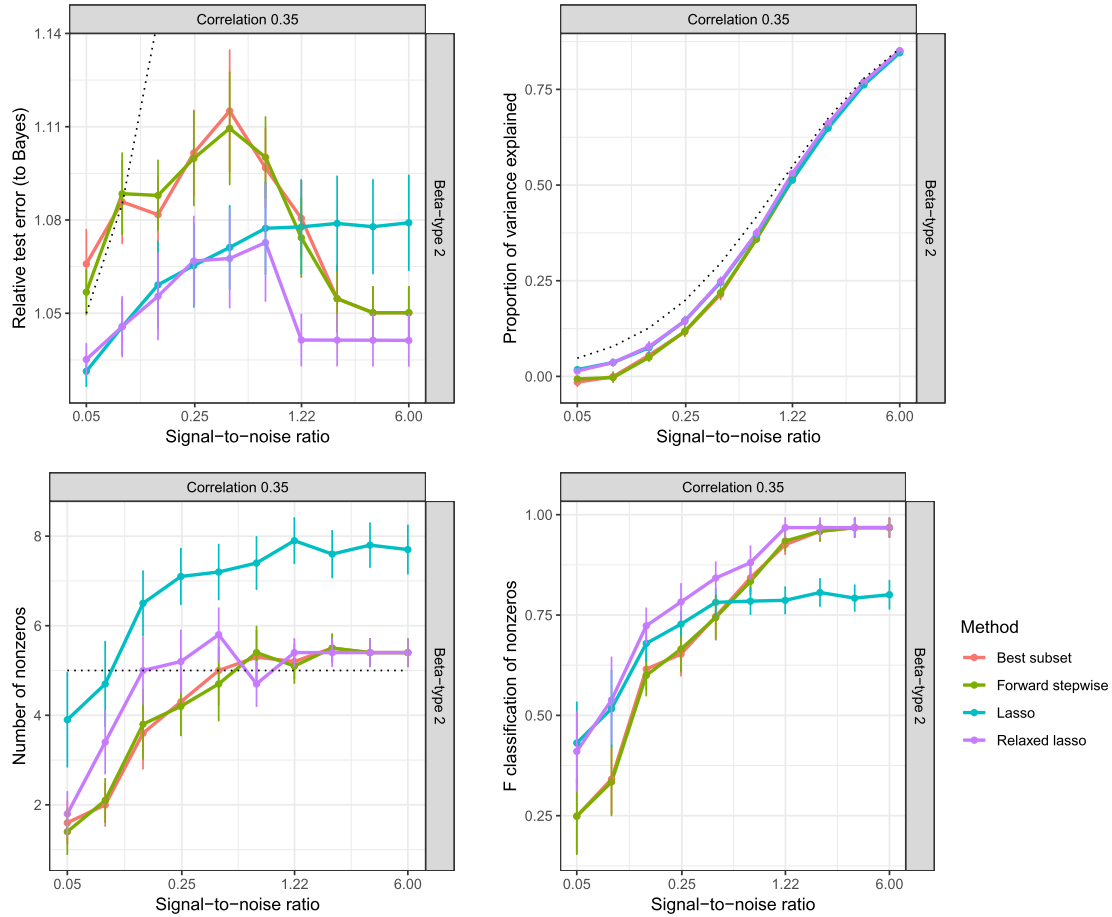**Correlation** $\rho = 0.35$, **beta-type 2**

Fig. 5. *RTE, PVE, number of nonzero coefficients and F-score as functions of SNR, in the low setting with* $n = 100$, $p = 10$ *and* $s = 5$.

Recall, the MIO solver for best subset was allowed 30 minutes per subset size k, or 1500 minutes for a path of 50 subset sizes. As the times in Table 1 suggest, the maximum allotted time was not reached in all instances, and the MIO solver managed to verify optimality of some solutions along the path. In the medium setting, on average 17.55 of the 50 solutions were verified as being optimal. In the high-5 and high-10 settings, only 1.61 of the 50 were verified on average (note this count includes the subset of size 1, which is trivial). These measures may be pessimistic, as Gurobi may have found high-quality approximate solutions or even exact solutions but was just not able to verify them in time; see the discussion in the above subsection.

3.4 Results: Accuracy Metrics

Here, we display a slice of the accuracy results, focusing for concreteness on the case in which the predictor correlation level is $\rho = 0.35$, and the population coefficients follow the beta-type 2 pattern. The only exception is in Figure 8, where we show the results for beta-type 1 setting. In a Supplementary Appendix (Hastie, Tibshirani

and Tibshirani, 2020), we display the full set of results, over the whole simulation design.

Figure 5 plots the relative test error, PVE, number of nonzero coefficients and F-score as functions of the SNR level, for the low setting. Figures 6 and 7, show the same for the medium and high-5 settings, respectively. Figure 8 shows the results for high-5 and a different parameter setting—beta-type 1, chosen because it represents the only scenario where best subset seems to show an advantage. Each panel in the figures shows the average of a given metric over 10 repetitions, for the four methods in question, and vertical bars denote one standard error. In the relative test error plots, the dotted curve denotes the performance of the null model (null score); in the PVE plots, it denotes the performance of the true model (perfect score); in the number of nonzero plots, it marks the true support size s.

The low and medium settings, Figures 5 and 6, yield somewhat similar results. In the relative test error and PVE plots (top left and top right panels), we see that best subset and forward stepwise lag behind the lasso and relaxed lasso in terms of accuracy for low SNR levels; as

Medium setting: $n = 500$, $p = 100$, $s = 5$
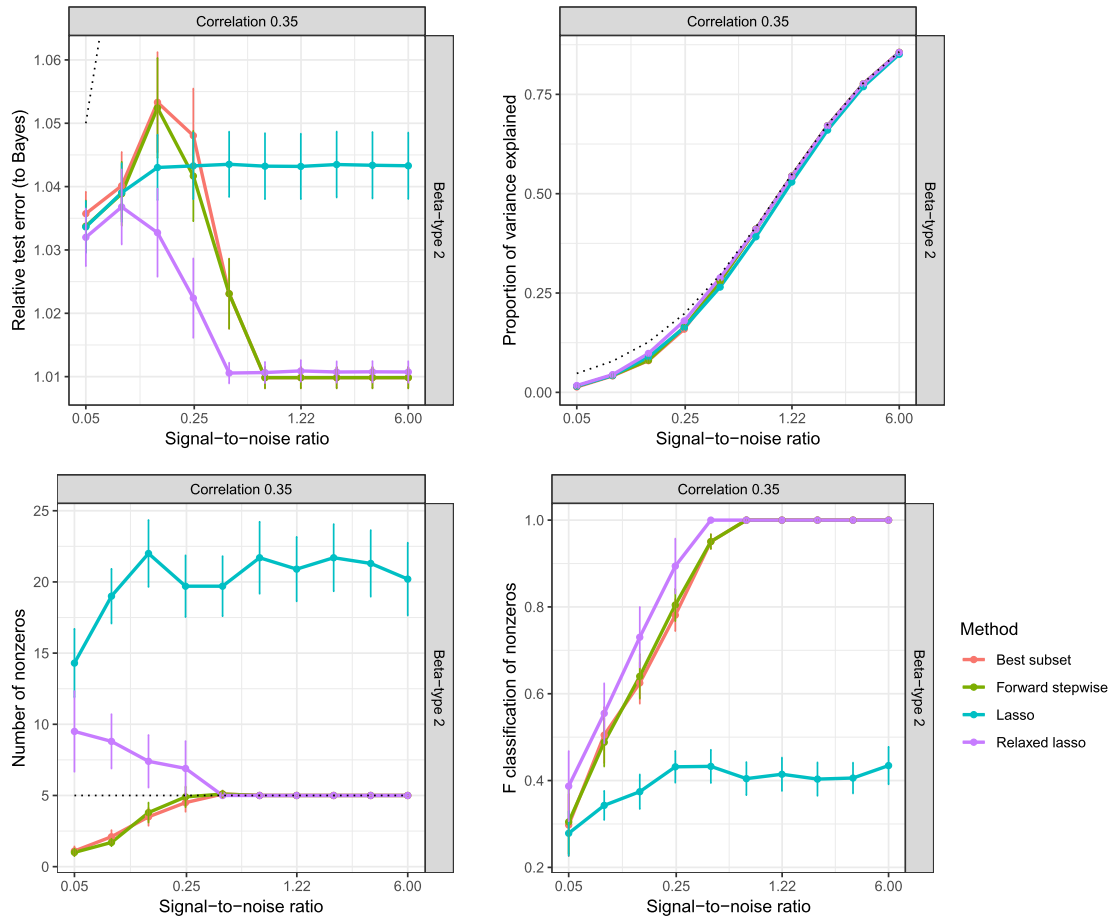Correlation $\rho = 0.35$, **beta-type 2**

FIG. 6. *RTE, PVE, number of nonzero coefficients and F-score as functions of SNR, in the medium setting with $n = 500$, $p = 100$ and $s = 5$.*

the SNR increases, we see in the PVE plot that all four methods converge to nearly perfect accuracy. The relative test error plot (top left panel) magnifies the differences between the methods. For low SNR levels, we see that the lasso outperforms the more aggressive best subset and forward stepwise methods, but for high SNR levels, it is outperformed by the latter two methods. The critical transition point—the SNR value at which their relative test error curves cross—is different for the low and medium settings: for the low setting, it is around 1.22, and for the medium setting, it is earlier, around 0.42. The relaxed lasso, meanwhile, is competitive across all SNR levels: at low SNR levels, it matches the performance of the lasso, and at high SNR levels, it matches that of best subset and forward stepwise. It is able to do so by properly tuning the amount of shrinkage (via its parameter γ) on the validation set. Lastly, the number of nonzero estimated coefficients from the four methods (bottom left panel) is also revealing. The lasso consistently delivers much denser solutions; essentially, to optimize prediction error on the validation set, it is forced to do so, as the

sparser solutions along its path entail too much shrinkage. The relaxed lasso does not suffer from this issue, again thanks to its ability to unshrink (move γ away from 1); it delivers solutions that are just as sparse as those from best subset and forward stepwise, except at the low SNR range.

In the high-5 setting of Figure 7, the methods behave quite differently. The PVEs delivered by all methods are close to zero for low SNR values. We see that there is no strong reason, based on relative test error, PVE or F-score, to favor best subset or forward stepwise over the lasso. However unlike the lasso, best subset and forward stepwise yield the correct null model in this setting, a desirable property. At low SNR levels, best subset and forward stepwise often have worse accuracy metrics (and certainly more erratic metrics); at high SNR levels, these procedures do not show much of an advantage. The relaxed lasso again performs the best overall, with a noticeable gap in performance at the high SNR levels. As is confirmed by the number of nonzero coefficients plots, the lasso and best subset/forward stepwise achieve sim-

High-5 setting: $n = 50, \ p = 1000, \ s = 5$
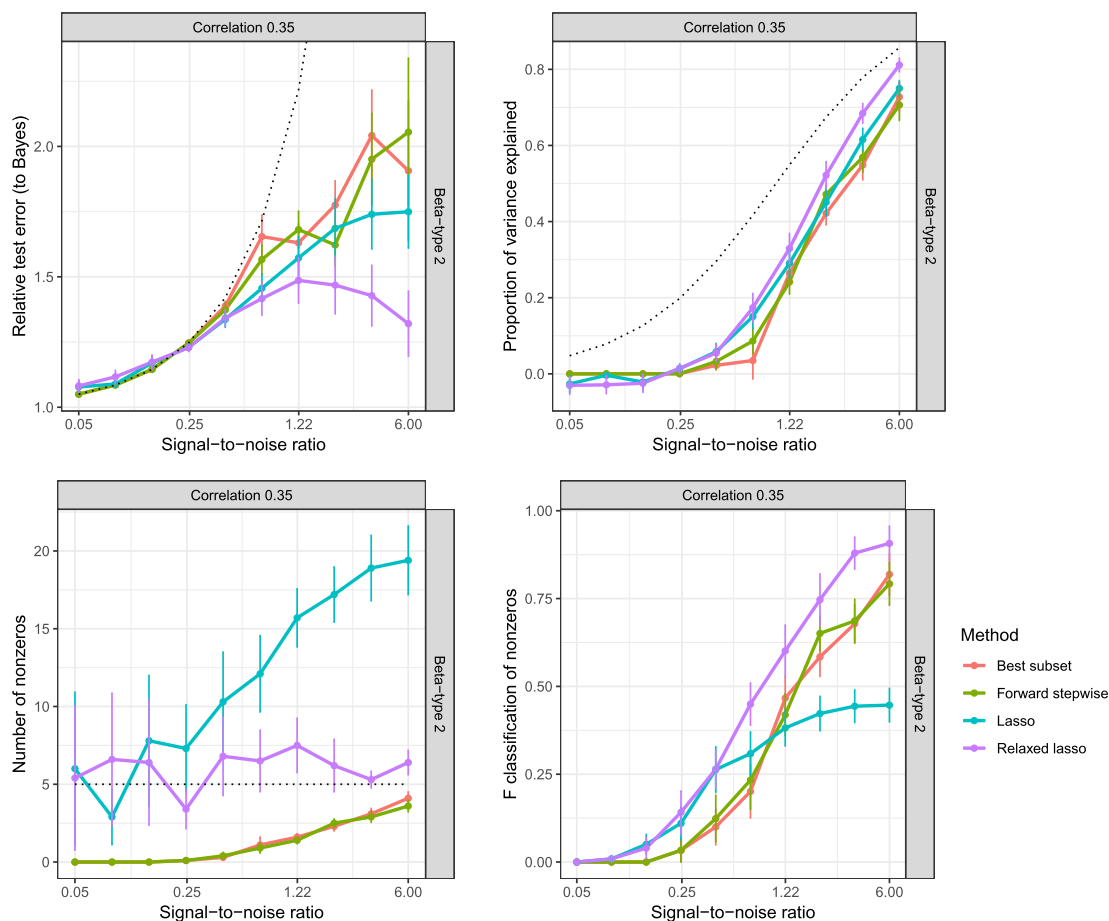**Correlation** $\rho = 0.35$, **beta-type 2**

FIG. 7.   *RTE, PVE, number of nonzero coefficients and F-score as functions of SNR, in the high-5 setting with $n = 50$, $p = 1000$ and $s = 5$.*

ilar accuracy in the high SNR range using two opposite strategies: the former uses high-bias and low-variance estimates, and the latter uses low-bias and high-variance estimates. The relaxed lasso is most accurate by striking a favorable balance between these two polar regimes.

We note however that there is one setting—Figure 8—where best subset is the winner (high-5, beta-type 1). When the SNR is above about 1.5, it shows clear gains over other methods. Recall that the beta type-1 setting has its non-zero coefficients on an equally-spaced grid from 1 to p. We postulate that the lasso (and relaxed lasso) have trouble in this case because the irrepresentability condition for the lasso does not hold (at least when the feature correlation is greater than zero). We note that this advantage for best subset does not carry over, empirically, to the high-10 setting. The details are given in the Supplementary Appendix (Hastie, Tibshirani and Tibshirani, 2020).

3.5 Summary of Results

As mentioned above, the results from our entire simulation suite can be found the Supplementary Material

(Hastie, Tibshirani and Tibshirani, 2020). Here is a high-level summary:

- Forward stepwise and best subset perform quite similarly throughout (with the former being much faster), although in some high SNR settings, best subset does perform better. This does not agree with the results for forward stepwise in Bertsimas, King and Mazumder (2016), where it performed quite poorly in comparison. In talking with the authors of that paper, we have learned that this was due to the fact that forward stepwise in their study was tuned using AIC, rather than a separate validation set. So, when put on equal footing and allowed to select its tuning parameter using validation data just as the other methods, we see that it performs quite comparably.

- The lasso gives better accuracy results than the best subset in the low SNR range and worse accuracy than the best subset in the high SNR range. The SNR transition point varies depending on the problem dimensions (n, p) predictor correlation level (ρ) and beta-type (1 through 5). For the medium setting, the tran-

High-5 setting: $n = 50$, $p = 1000$, $s = 5$
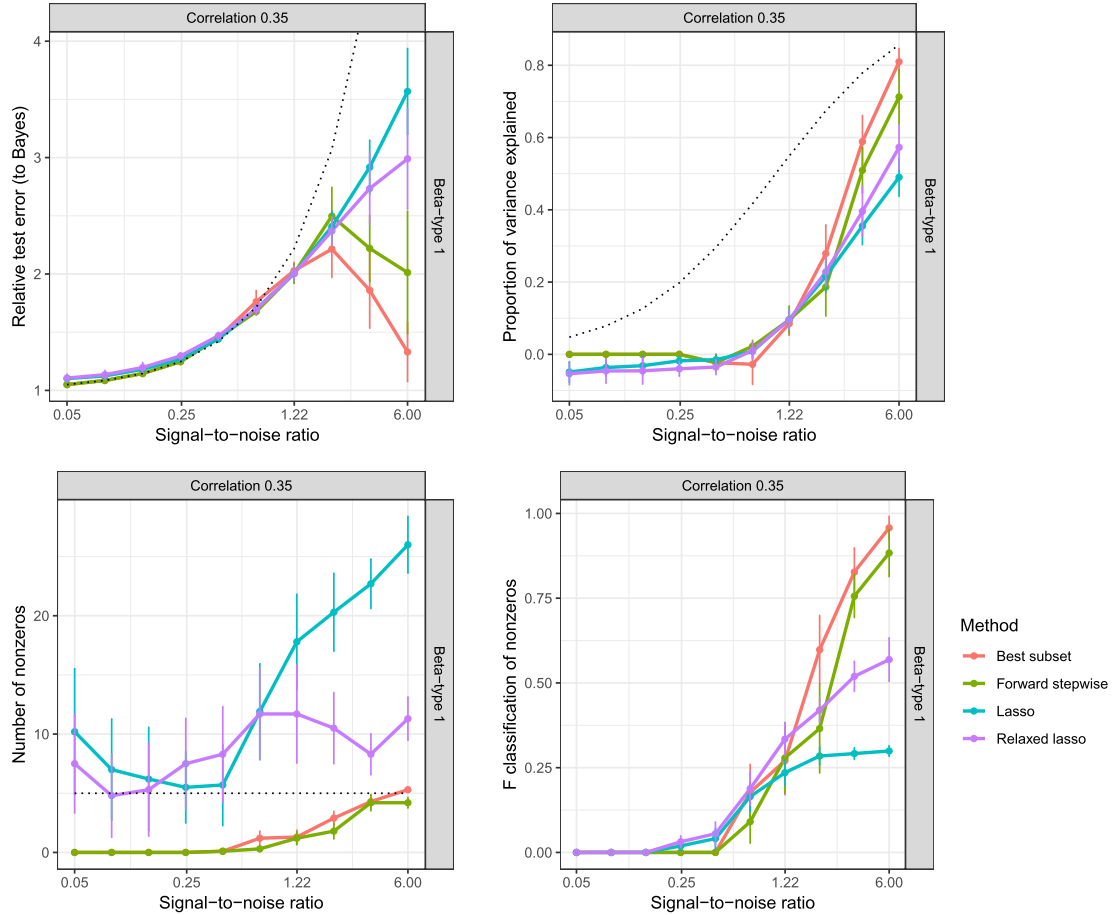Correlation $\rho = 0.35$, beta-type 1

Fig. 8. *RTE, PVE, number of nonzero coefficients and F-score as functions of SNR, in the high-5 setting with $n = 50$, $p = 1000$ and $s = 5$.*

sition point comes earlier than in the low setting. For the high-10 setting, the transition point often does not come at all (before an SNR of 6, which is the maximum value we considered). For the high-5 setting, this transition only seems to occur for the beta type-1 model. As the predictor correlation level increases, the transition point typically appears later (again, in some cases it does not come at all, e.g., for beta-type 5 and correlation level $\rho = 0.7$).

- The relaxed lasso provides overall the top accuracy results. In nearly all cases (across all SNR levels, and in all problem configurations), we considered, it performs as well as or better than all other methods. We conclude that it is able to use its auxiliary shrinkage parameter (γ) to get the "best of both worlds": it accepts the heavy shrinkage from the lasso when such shrinkage is helpful, and reverses it when it is not.
- The PVE plots remind us that, despite what may seem like large relative differences, the four methods under consideration do not have very different absolute performances in this intuitive and important metric. It thus

makes sense overall to favor the methods that are easy to compute.
- The Gurobi MIO algorithm for best subset was given 30 minutes per problem (i.e., for each subset size). For a typical grid of 50 subset sizes, this results in around 25 hours of computation. When used within 10-fold CV, this increases to 250 hours or more than 10 days of computation. By comparison, all three other methods take seconds.

4. DISCUSSION

The recent work of Bertsimas, King and Mazumder (2016) has enabled the first large-scale empirical examinations of best subset. In this paper, we have expanded and refined the simulations in their work, comparing best subset to forward stepwise, the lasso and the relaxed lasso. We have found: (a) forward stepwise and best subset perform similarly throughout, with a few exceptions in the high SNR scenario; (b) best subset often loses to the lasso except in the high SNR range; (c) the relaxed lasso achieves "the best of both worlds" and per-

forms on par with the best method in almost every scenario.

Our R package bestsubset, designed to easily replicate all of the simulations in this work, or forge new comparisons, is available at https://github.com/ryantibs/best-subset/.

ACKNOWLEDGMENTS

We thank Rahul Mazumder for his help and guidance, and Balasubramanian Narasimhan for his help with the computing setup.

SUPPLEMENTARY MATERIAL

Supplement to "Best Subset, Forward Stepwise or Lasso? Analysis and Recommendations Based on Extensive Comparisons" (DOI: 10.1214/19-STS733SUPP; .pdf). The supplementary material provides the complete set of results from our full simulation suite.

REFERENCES

BEALE, E. M. L., KENDALL, M. G. and MANN, D. W. (1967). The discarding of variables in multivariate analysis. *Biometrika* **54** 357–366. MR0221656 https://doi.org/10.1093/biomet/54.3-4.357

BELLONI, A., CHERNOZHUKOV, V. and WANG, L. (2011). Square-root lasso: Pivotal recovery of sparse signals via conic programming. *Biometrika* **98** 791–806. MR2860324 https://doi.org/10.1093/biomet/asr043

BERTSIMAS, D., KING, A. and MAZUMDER, R. (2016). Best subset selection via a modern optimization lens. *Ann. Statist.* **44** 813–852. MR3476618 https://doi.org/10.1214/15-AOS1388

BERTSIMAS, D. and VAN PARYS, B. (2020). Sparse high-dimensional regression: Exact scalable algorithms and phase transitions. *Ann. Statist.* **48** 300–323. MR4065163 https://doi.org/10.1214/18-AOS1804

CANDES, E. and TAO, T. (2007). The Dantzig selector: Statistical estimation when p is much larger than n. *Ann. Statist.* **35** 2313–2351. MR2382644 https://doi.org/10.1214/009053606000001523

CHEN, S. S., DONOHO, D. L. and SAUNDERS, M. A. (1998). Atomic decomposition by basis pursuit. *SIAM J. Sci. Comput.* **20** 33–61. MR1639094 https://doi.org/10.1137/S1064827596304010

DAVIS, G., MALLAT, S. and ZHANG, Z. (1994). Adaptive time-frequency decompositions with matching pursuit. *Wavelet Appl.* **402** 402–413.

DRAPER, N. R. and SMITH, H. (1966). *Applied Regression Analysis*. Wiley, New York. MR0211537

EFRON, B., HASTIE, T., JOHNSTONE, I. and TIBSHIRANI, R. (2004). Least angle regression. *Ann. Statist.* **32** 407–499. MR2060166 https://doi.org/10.1214/009053604000000067

EFROYMSON, M. (1966). Stepwise regression—A backward and forward look. Paper presented at the *Eastern Regional Meetings of the Institute of Mathematical Statistics*.

FAN, J. and LI, R. (2001). Variable selection via nonconcave penalized likelihood and its oracle properties. *J. Amer. Statist. Assoc.* **96** 1348–1360. MR1946581 https://doi.org/10.1198/016214501753382273

FRIEDMAN, J., HASTIE, T. and TIBSHIRANI, R. (2010). Regularization paths for generalized linear models via coordinate descent. *J. Stat. Softw.* **33** Art. ID 1.

FRIEDMAN, J., HASTIE, T., HÖFLING, H. and TIBSHIRANI, R. (2007). Pathwise coordinate optimization. *Ann. Appl. Stat.* **1** 302–332. MR2415737 https://doi.org/10.1214/07-AOAS131

FURNIVAL, G. and WILSON, R. (1974). Regression by leaps and bounds. *Technometrics* **16** 499–511.

GOLUB, G. H. and VAN LOAN, C. F. (1996). *Matrix Computations*, 3rd ed. *Johns Hopkins Studies in the Mathematical Sciences*. Johns Hopkins Univ. Press, Baltimore, MD. MR1417720

HASTIE, T., TIBSHIRANI, R. and FRIEDMAN, J. (2009). *The Elements of Statistical Learning*: *Data Mining*, *Inference*, *and Prediction*, 2nd ed. *Springer Series in Statistics*. Springer, New York. MR2722294 https://doi.org/10.1007/978-0-387-84858-7

HASTIE, T., TIBSHIRANI, R. and TIBSHIRANI, R. (2020). Supplement to "Best Subset, Forward Stepwise or Lasso? Analysis and Recommendations Based on Extensive Comparisons." https://doi.org/10.1214/19-STS733SUPP

HAZIMEH, H. and MAZUMDER, R. (2018). Fast best subset selection: Coordinate descent and local combinatorial optimization algorithms. Preprints. Available at arXiv:1803.01454.

HOCKING, R. R. and LESLIE, R. N. (1967). Selection of the best subset in regression analysis. *Technometrics* **9** 531–540. MR0219192 https://doi.org/10.2307/1266192

JANSON, L., FITHIAN, W. and HASTIE, T. J. (2015). Effective degrees of freedom: A flawed metaphor. *Biometrika* **102** 479–485. MR3371017 https://doi.org/10.1093/biomet/asv019

KAUFMAN, S. and ROSSET, S. (2014). When does more regularization imply fewer degrees of freedom? Sufficient conditions and counterexamples. *Biometrika* **101** 771–784. MR3286916 https://doi.org/10.1093/biomet/asu034

MALLAT, S. and ZHANG, Z. (1993). Matching pursuits with time-frequency dictionaries. *IEEE Trans. Signal Process.* **41** 3397–3415.

MAZUMDER, R., FRIEDMAN, J. H. and HASTIE, T. (2011). *SparseNet*: Coordinate descent with nonconvex penalties. *J. Amer. Statist. Assoc.* **106** 1125–1138. MR2894769 https://doi.org/10.1198/jasa.2011.tm09738

MEINSHAUSEN, N. (2007). Relaxed Lasso. *Comput. Statist. Data Anal.* **52** 374–393. MR2409990 https://doi.org/10.1016/j.csda.2006.12.019

NATARAJAN, B. K. (1995). Sparse approximate solutions to linear systems. *SIAM J. Comput.* **24** 227–234. MR1320206 https://doi.org/10.1137/S0097539792240406

RADCHENKO, P. and JAMES, G. M. (2011). Improved variable selection with forward-Lasso adaptive shrinkage. *Ann. Appl. Stat.* **5** 427–448. MR2810404 https://doi.org/10.1214/10-AOAS375

TIBSHIRANI, R. (1996). Regression shrinkage and selection via the lasso. *J. Roy. Statist. Soc. Ser. B* **58** 267–288. MR1379242

TIBSHIRANI, R. J. (2013). The lasso problem and uniqueness. *Electron. J. Stat.* **7** 1456–1490. MR3066375 https://doi.org/10.1214/13-EJS815

TIBSHIRANI, R. J. and TAYLOR, J. (2012). Degrees of freedom in lasso problems. *Ann. Statist.* **40** 1198–1232. MR2985948 https://doi.org/10.1214/12-AOS1003

TIBSHIRANI, R., BIEN, J., FRIEDMAN, J., HASTIE, T., SIMON, N., TAYLOR, J. and TIBSHIRANI, R. J. (2012). Strong rules for discarding predictors in lasso-type problems. *J. R. Stat. Soc. Ser. B. Stat. Methodol.* **74** 245–266. MR2899862 https://doi.org/10.1111/j.1467-9868.2011.01004.x

ZHANG, C.-H. (2010). Nearly unbiased variable selection under minimax concave penalty. *Ann. Statist.* **38** 894–942. MR2604701 https://doi.org/10.1214/09-AOS729

ZOU, H., HASTIE, T. and TIBSHIRANI, R. (2007). On the "degrees of freedom" of the lasso. *Ann. Statist.* **35** 2173–2192. MR2363967 https://doi.org/10.1214/009053607000000127